

Workshop 1: Introduction to scikit-learn

TECH6

Richard Foltyn
NHH Norwegian School of Economics

October 15, 2026

See GitHub repository for notebooks and data:

<https://github.com/richardfoltyn/TECH6-H26>

Exercise 1: Polynomial under- and overfitting

Consider the following non-linear model,

$$y_i = f(x_i) + \epsilon_i$$
$$f(x) = \cos\left(\frac{3}{2}\pi x\right)$$
$$\epsilon_i \stackrel{\text{iid}}{\sim} N(0, \sigma_\epsilon^2)$$

where y_i is a trigonometric function of x_i but is measured with an additive error ϵ_i . In this exercise, we are going to approximate y_i using polynomials in x_i of varying degrees.

Part 1 — Creating and plotting a sample

The function `fcn()` implements the true relationship $y = f(x)$.

```
[1]: import numpy as np

def fcn(x):
    """
    True function without errors.
    """
    return np.cos(1.5 * np.pi * x)
```

Using this function, create a sample of size $N = 50$ where the x_i are randomly drawn from a uniform distribution on the interval $[0, 1]$ and $\sigma_\epsilon = 0.2$ (initialize your RNG with a seed of 1234). Then generate y_i according to the equation given above.

Create a scatter plot of the sample (x_i, y_i) and add a line depicting the true relationship without measurement error.

Part 2 — Polynomial approximations

In order to approximate this relationship using polynomials, you first want to visualize how the polynomial degree affects the model fit. To this end, complete the template for `plot_poly_degrees()` below to create a figure with 6 subplots, each showing the predicted values for a polynomial of degree d , where

$d \in \{0, 1, 2, 3, 10, 15\}$. Each panel should additionally show the sample scatter plot and the true function $y = f(x)$.

How does the quality of the approximation change as you increase d ? Do higher-order polynomials always perform better?

Hints:

- You should build a pipeline, e.g., using `make_pipeline()` which combines the `PolynomialFeatures` transformation and `LinearRegression` to approximate y as a polynomial in x .
- When creating polynomials with `PolynomialFeatures(..., include_bias=True)` you need to fit the model *without* an additional intercept as the intercept is already included in the polynomial.

```
[4]: import matplotlib.pyplot as plt

def plot_poly_degrees(degrees, X, y):
    """
    Fit and plot polynomial regression models of different degrees.

    Parameters
    -----
    degrees : array-like
        List of polynomial degrees to fit and plot.
    X : array-like
        Explanatory variable.
    y : array-like
        Response variable.
    """

    # Create figure with 3 columns
    ncol = 3
    nrow = int(np.ceil(len(degrees) / ncol))

    fig, axes = plt.subplots(
        nrow, ncol, figsize=(8, 5), sharex=True, sharey=True, constrained_layout=True
    )

    # x-values for predicting & plotting
    xvalues = np.linspace(0.0, 1.0, 101)
    # True y-values
    y_true = fcn(xvalues)

    for i, ax in enumerate(axes.flatten()):
        # Degree for current sub-plot
        d = degrees[i]

        # TODO: Create polynomial transformation + estimation pipeline

        # TODO: Fit the model to the sample data

        # TODO: Predict the response variable for the x-values

        # TODO: Plot predicted values

        # Plot true relationship
        ax.plot(xvalues, y_true, color='black', lw=1.0)
        # Plot sample as scatter plot
        ax.scatter(X, y, s=20, color='none', edgecolor='steelblue', lw=0.75)
        ax.set_ylim((-1.5, 1.5))
        ax.legend()
```

Part 3 — Compute MSE for each hyperparameter

You want to find the optimal polynomial degree using cross-validation. To this end, implement the function `compute_average_mse()` using the template below.

This function takes as arguments the polynomial degree d , the sample observations (X, y) , and the number of splits `n_splits`, and returns the mean squared error (MSE) for the training and validation samples averaged across all splits.

Hint: To use the MSE as the scoring metric in `cross_validate()`, pass the argument `scoring='neg_mean_squared_error'` and remember to negate the output to get the actual MSE.

```
[7]: from sklearn.model_selection import cross_validate

def compute_average_mse(d, X, y, n_splits=10):
    """
    Compute mean squared error averaged across splits in k-fold cross-validation.

    Parameters
    -----
    d : int
        Degree of polynomial.
    X : array-like
        Explanatory variable.
    y : array-like
        Response variable.
    n_splits : int
        Number of splits in k-fold cross-validation.

    Returns
    -----
    mse_train : float
        MSE on training sample averaged over all splits.
    mse_validate : float
        MSE on validation sample averaged over all splits.
    """

    # TODO: Create pipeline object

    # TODO: Run cross-validation using cross_validate()

    # TODO: Compute MSE on train and validation samples averaged over all splits

    # TODO: Return average MSE on train and validation samples
```

Part 4 — Optimal polynomial degree

Using the function you wrote, compute the average MSEs for a sequence of 16 polynomial degrees $d = 0, 1, 2, \dots, 15$ using 10 splits. Use the MSE statistics to plot the validation curve showing the average MSE on the y -axis against d on the x -axis. Which degree d results in the lowest average MSE on the validation sample?

Part 5 — Plotting the fitted model

Re-estimate the model using the optimal polynomial degree you just found and create a scatter plot with the original data, the true function $y = f(x)$, and the fitted polynomial.

Part 6 — Automatic cross-validation

The steps in Parts 3 and 4 can be implemented more easily using `validation_curve()`. Re-implement the cross-validation using this function by completing the code template below.

Hint: Don't forget that you have to use the *negative* MSE as the relevant criterion, i.e., specify the argument `scoring='neg_mean_squared_error'` when calling `validation_curve()`.

```
[14]: from sklearn.model_selection import validation_curve

# TODO: Create estimation pipeline

# TODO: Complete the call to validation_curve()
# train_scores, validate_scores = validation_curve(
#     estimator=...,
#     X=X, y=y,
#     param_name=...,
#     param_range=...,
#     scoring=...,
#     cv=10
# )

# TODO: Average train_scores and validate_scores across CV splits

# TODO: Report the polynomial degree that minimizes the MSE on the validation sample
```

Exercise 2: Categorical features

In this exercise, we practice fitting models that include categorical features in the regressor matrix X using the Ames housing dataset. There are several categorical variables in this data that come in two different kinds:

1. **Nominal categorical variables** capture qualitative categories. For numerical purposes, we often encode these as integer values, but any particular integer representation is arbitrary. It is a serious mistake to include such integer-valued categorical variables as *continuous* features in a model.

In the Ames housing data, the Neighborhood column is one such nominal categorical variable (stored as a string containing the neighborhood name).

2. **Ordinal categorical variables** represent data that have a natural ordering, but the distance between any two values can be arbitrary.

For example, in the Ames housing data the OverallQuality variable is an ordinal categorical variable on a scale of 1–10. While a value of 2 is clearly better than a value of 1, it is up to interpretation how much better.

Such variables can be included as categorical *dummy variables* in a model, not imposing any particular interpretation on the distance between two values. Alternatively, such variables can be included as *continuous* variables, which imposes how the distance is interpreted. For example, including an ordinal categorical variable as a linear term in a model imposes that moving from 1 to 2 has the same effect on the outcome variable as moving from 2 to 3.

Part 1 — Creating the estimation sample

We load the Ames data the same way as we did in the lectures:

```
[19]: import pandas as pd

# Use local data/ folder
DATA_PATH = '../data'
```

```
filename = f'{DATA_PATH}/ames_houses.csv'
df = pd.read_csv(filename, sep=',')
```

First, perform the following data preprocessing steps:

1. Drop all observations that have any missing values for SalePrice, OverallQuality, or Neighborhood.
2. Tabulate how many neighborhoods have fewer than 40 observations, and drop observations from these neighborhoods from the dataset.

Part 2 — Exploratory data analysis

Before fitting a model, you should perform some exploratory data analysis to familiarize yourself with the data:

1. Compute the correlations of SalePrice with all numerical variables in the data. Which variable is the most correlated?
2. Create a histogram for the variable OverallQuality, showing the number of observations for each of the possible values 1–10.

Hint: Create a Series with the number of observations per quality level and plot it as a bar chart using `Series.plot.bar()`.

3. Create a box plot showing the SalePrice by OverallQuality.

Hint: You can use the pandas plotting function `DataFrame.plot.box()` for this.

4. Create a box plot showing the SalePrice by Neighborhood.

Part 3 — Train-test split

Split the data into a training and a test sample, assigning 40% of observations to the test sample, and report the number of observations in each set.

Stratify your split by Neighborhood so that each neighborhood is approximately equally represented in the training and test samples (use the `stratify` argument of `train_test_split()` for this). Specify a random seed of 1234.

Note that you can pass the whole DataFrame to `train_test_split()` instead of individual X and y values. The function will then return two DataFrames, one containing the training and one containing the test data.

Part 4 — Predict sale price by neighborhood

Create an estimator pipeline to fit the model

$$\text{SalePrice}_{ij} = \mu_j + \epsilon_{ij}$$

where i indexes observations and j indexes neighborhoods. That is, the sale price for a house i in neighborhood j is predicted to be the average sale price μ_j in that neighborhood.

In order to estimate this model, you'll have to convert the categorical string variable Neighborhood to a set of 0/1 dummy variables. This can be achieved using the `OneHotEncoder` as a step in a pipeline.

Fit the model on the training data, and report the RMSE both on the training and the test sample.

You are now asked to fit the following model:

$$\text{SalePrice}_i = \beta_0 + \beta_1 \text{OverallQuality} + \beta_2 \text{OverallQuality}^2 + \dots + \beta_K \text{OverallQuality}^K + \epsilon_i$$

where the sale price is assumed to be a function of a polynomial in overall quality where we vary the polynomial degree K . We are thus implicitly treating the ordinal categorical variable OverallQuality as a continuous variable, which may or may not be admissible depending on the data. We will compare the estimation results to treating OverallQuality as a proper categorical variable in the next part.

Use the following template code to fit the above model for polynomial degrees $d = 1, 2, \dots, 9$ and compute the cross-validated RMSE on the training and validation samples. For this you should use the function `cross_validate()`. This function returns a dictionary with the keys 'train_score' and 'test_score' that contain the computed scores for each split. Note that you need to pass `return_train_score=True`, otherwise the function does not return the scores for the training data.

Create a plot showing the validation curves for both the training and validation samples, and report the polynomial degree that minimizes the RMSE on the validation sample.

```
[34]: import numpy as np

degrees = np.arange(1, 10)

rmse_train_mean = []
rmse_validate_mean = []

for d in degrees:
    # TODO: Create polynomial features + estimation pipeline

    # TODO: Compute RMSE using cross-validation
    # scores = cross_validate(
    #     estimator=...,
    #     X=X_train, y=y_train,
    #     scoring=...,
    #     cv=5,
    #     return_train_score=True,
    #     n_jobs=-1
    # )

    # TODO: Compute average RMSE over all CV splits
    # rmse_train_mean.append(np.mean(-scores['train_score']))
    # rmse_validate_mean.append(np.mean(-scores['test_score']))

pass
```

Part 6 — Predict price with quality dummies

You are now interested in whether treating the overall quality as a proper categorical variable improves the model performance. To this end, estimate the model

$$\text{SalePrice}_{ij} = \mu_j + \epsilon_{ij}$$

where i indexes observations and j indexes the overall quality categories 1–10. That is, the sale price for a house i with overall quality j is predicted to be the average sale price μ_j of all houses with that overall quality.

Create an estimation pipeline to fit the model and perform cross-validation exactly as you did for the polynomial case above. Note that cross-validation is not needed for this model at all since it has no hyperparameters, but doing so creates a fair comparison to the polynomial case since we are estimating and predicting on the same sample using the same splits.

Recreate the validation curve from above and add two horizontal lines showing the average RMSE from the dummy variable model for the training and validation sets.

```
[39]: # TODO: Create pipeline with dummy variables + estimation

# TODO: Run cross-validation
```

```
# scores = cross_validate(  
#     estimator=...,  
#     X=X_train, y=y_train,  
#     scoring=...,  
#     cv=5,  
#     return_train_score=True,  
#     n_jobs=-1  
# )  
  
# TODO: Compute average RMSE over all CV splits  
# rmse_train_dummies_mean = np.mean(-scores['train_score'])  
# rmse_validate_dummies_mean = np.mean(-scores['test_score'])
```

Part 7 — Compare test sample predictions

Finally, use the best polynomial model and the dummy variable model estimated on the whole training sample, and compute the RMSE on the test sample. Which model performs better?